



Daniel Fridljand

📍 Munich, Germany ✉ danielfridljand.de in daniel-fridljand 🎧 FridljDa

Summary

Software consultant at TNG focused on applied AI engineering, currently embedded as the sole AI engineer on a production agentic support-automation system (Claude, Temporal, MCP, DSPy) at an enterprise customer. Research background spanning ETH Zürich (computational oncology), Stanford School of Medicine (first co-author in *Nature Medicine*, 2024), and EMBL (statistical genomics). M.Sc. Mathematics with full marks; Yale exchange scholar.

Software Development Experience

Software Consultant - Applied AI

TNG Technology Consulting

Munich, Germany

Dec 2025 – present

- **Secure Desktop AI Agent (Jun 2026 – present)**: Core developer on a privacy-first desktop AI agent (Tauri/React frontend, Rust backend); full-stack contributions across the desktop UI, agent runtime, and secure auditable LLM tool execution.
- **LLM-Powered Document Validation System (May–Jun 2026)**: Core developer in a 3-person team for a production-grade automated proposal review service with .docx processing and automated comment insertion.
- **Email-to-Order Parsing (Apr–May 2026)**: Designed a field-centric email-processing pipeline with FastAPI ingestion, Streamlit review dashboard, and attachment-aware multimodal LLM processing (vision fallbacks for scanned PDFs) with parallel candidate resolution.
- **Customer Support Automation (Dec 2025 – Apr 2026, live in production)**: Sole AI engineer end-to-end on a cinema-ticketing SaaS — **770+ live B2C tickets at 86.6% strict / 88.9% content-supported approval** daily volume scaled 5x after CEO showcase; 13+ feature improvements from live reviewer feedback.
- Architected a hybrid deterministic + agentic resolution workflow over 19 customer-intent categories using **Temporal** for durable orchestration with human-in-the-loop approval via Signals, and **DSPy/GEPA** for prompt optimization.
- Built a production evaluation suite in **Langfuse** (6 evaluation types, 8 score metrics, 14-label outcome taxonomy, LLM-as-judge) and a reproducible 22-stage **Snakemake** pipeline with GDPR-compliant Microsoft Presidio PII detection.
- Tech: Python, FastAPI, Streamlit, PydanticAI, Temporal, Langfuse, Snakemake, Docker, DSPy, MCP, Tauri, React, Rust.

Software Consultant - Enterprise Modernization

TNG Technology Consulting

Munich, Germany

Dec 2024 – Dec 2025

- Modernized a mission-critical global supply-chain application (Java 8 → 17, JBoss → WildFly) as part of the platform team in a multi-year transformation program.
- Shipped a JFrog Artifactory proxy in 3 days, reducing a recurring CI pipeline runtime from 8 hours to 30 seconds — saving developers ~1–2 hours per week each.
- Established DevSecOps governance: OWASP scans in CI, Grafana CVE dashboards, and SOAP→REST migration with Keycloak; architected daemonless Docker/Podman build environments.

Data Science Experience

Research Data Analyst - Computational Oncology

ETH Zürich - Department of Biosystems Science and Engineering

Basel, Switzerland

Feb 2024 – Sept 2024

- Developed novel Bayesian non-parametric methods (Hierarchical Dirichlet Process) for mutational signature estimation, extended to incorporate phylogenetic tree structures.
- Analyzed single-cell whole-exome sequencing data from the Tumor Profiler Study (187 cells, 10 melanoma tumors), identifying eight latent mutational signatures.
- Implemented hierarchical dependency structures in R where signature distributions for child phylogenetic nodes are drawn from parent distributions, enforcing biological inheritance.

Research Data Analyst - Environmental Epidemiology

Palo Alto, USA

- **Nature Medicine (2024)**: Led the statistical analysis as first co-author, quantifying air pollution's contribution to racial and socioeconomic mortality disparities in the US.
- Engineered a big-data pipeline harmonizing 63+ million death records, satellite pollution estimates, and census demographics across 3,000+ US counties (1990–2016).
- Implemented confounder-adjusted causal inference (DAGs, propensity scoring), revealing >50% of the Black–White all-cause mortality difference is attributable to environmental factors.
- Built an R Shiny analytical application used directly by epidemiologists and policy researchers to explore 17-dimensional data.

Research Data Analyst - Statistical Genomics (Master's Thesis)

Heidelberg, Germany

European Molecular Biology Laboratory (EMBL)

Oct 2021 – May 2022

- Developed **IHW-Forest**, a novel multiple-testing method using Random Forests for hypothesis weighting, increasing discovery power by >30% on 16 billion genetic association tests.
- Optimized core splitting and weighting logic in C++ via Rcpp for high-performance processing of large-scale genomic data.
- Presented at seven scientific events including Yale University and a competitively selected oral contribution at DAGStat 2022.

Education**University of Heidelberg**

Heidelberg, Germany

M.Sc. in M.Sc. Mathematics (Top of Class)

Oct 2020 – May 2023

- Grade: 1.0 (summa cum laude). Focus: ML, High-Dimensional Statistics, Probability Theory.
- Yale Exchange Scholar (2022–23, Honors): Deep Learning; Geometric & Topological ML (Prof. Krishnaswamy); Differentiable Manifolds.
- Master's Thesis (Grade 1.0) at EMBL: 'Better multiple Testing: Using multivariate co-data for hypothesis weighting'.

Publications**Disparities in air pollution attributable mortality in the US population by race, ethnicity and sociodemographic factors**

July 2024

Pascal Geldsetzer, ***Daniel Fridljand***, Mathew V. Kiang, Eran Bendavid, Sam Heft-Neal, Marshall Burke, et al.
www.nature.com/articles/s41591-024-03117-0 (Nature Medicine)

Projects**Data Mining Hackathon - Core Demand Prediction (TUM.ai x Unite)**

Online

Mar 2026

- **3rd in Level 2** (net €423K) and **5th in Level 1** (net €1.33M) with a two-stage LightGBM model (recurrence classifier + spend regressor) combined via an expected-utility objective.
- Engineered 200+ features and a reproducible Snakemake pipeline with automated portal submission via Playwright.

Agent Olympics Hackathon Munich 2026

Munich

Jan 2026

- Engineered a white-hat voice agent executing automated prompt-injection attacks across email, SMS, and WhatsApp, with automated OSINT-driven persona profiling.

Skills

Programming: Python (pandas, FastAPI, PyTorch, PydanticAI, DSPy), R (tidyverse, Shiny, Rcpp, Bioconductor), Java 17 (Spring), TypeScript/React, C++, SQL.

Applied AI: Agentic systems, LLM orchestration (PydanticAI, DSPy, MCP), LLM evaluation (Langfuse), prompt engineering, Causal Inference, Bayesian Non-parametrics.

Infrastructure: Docker, Podman, Kubernetes, CI/CD (Jenkins, GitHub Actions), AWS, Temporal, Snakemake, Prometheus, Grafana.

Languages: English (C2), German (Native), Russian (Native).