



Daniel Fridljand

danielfridljand.de
 daniel.fridljand@gmail.com
 daniel-fridljand
 FridljDa

Summary

AI engineer building production systems in Python and Rust, with a mathematician's rigor about whether they work. My proudest work to date: an LLM-based customer-support system I took from discovery to live production as the sole AI engineer, drafting resolutions for 800+ B2C tickets in its first three months at ~87% reviewer approval.

Industry Experience

AI Engineer - Software Consultant

TNG Technology Consulting

Munich, Germany

Dec 2025 – present

Building AI and predictive-analytics systems for clients and for TNG, from discovery through live operation.

- Architected a cinema-ticketing support system as a hybrid: deterministic paths for the recurring-intent majority, agentic fallback for the long tail, every draft reviewed by staff before customer contact.
- Built a feedback-driven improvement loop: reviewer corrections became shipped code changes, each gated against a growing Langfuse evaluation suite, on Temporal workflows that pause for staff sign-off.
- Sole engineer on a scrap-prediction proof of concept for a manufacturer: across ~40 benchmarked models, order attributes fixed at booking matched the full sensor set, steering the client toward tracking raw-material batches.
- One of three developers on an internal sandboxed platform (Tauri/React, Rust) for agentic workflows on sensitive local files; it cleared a 15-person adversarial red-team, unlocking expansion beyond the pilot.
- Scaled a retailer's cashier-fraud anomaly scoring from a 7-store sample to its full estate via SQL Server aggregation.

Platform Engineer - Software Consultant

TNG Technology Consulting

Munich, Germany

Dec 2024 – Dec 2025

Platform engineer on a multi-team program modernizing a global supply-chain application from Java 8 to 17.

- Deployed a JFrog Artifactory proxy in three days after the work had been repeatedly postponed as too costly, cutting a recurring CI step from eight hours to 30 seconds — saving each developer ~1–2 hours a week.
- Added OWASP dependency scanning to the Jenkins pipeline and built Grafana dashboards over Prometheus metrics, giving stakeholders visibility of CVE reduction throughout the migration.

Research Experience

Research Data Analyst - Environmental Epidemiology

Stanford University School of Medicine

Palo Alto, USA

July 2023 – Dec 2023

Statistical analysis and data tooling for a national study of air pollution and mortality inequality.

- Led the statistical analysis as co-first author of a [Nature Medicine paper](#) (2024), which found that over half the Black–White age-adjusted mortality gap in the US is attributable to fine-particulate air pollution.
- Built a reproducible pipeline integrating 63M+ death records with pollution estimates and demographic data; open-sourced the analysis code and interactive R Shiny explorer.

Research Data Analyst - Statistical Genomics (Master's Thesis)

European Molecular Biology Laboratory (EMBL)

Heidelberg, Germany

Oct 2021 – May 2022

Research on multiple-testing methods for large-scale genomics.

- Developed IHW-Forest, a random-forest hypothesis-weighting method with its splitting and weighting logic written in C++ via Rcpp — it recovered over 30% more discoveries than standard corrections.

Education

University of Heidelberg

M.Sc. in Mathematics

Heidelberg, Germany

Oct 2020 – May 2023

- Grade: 1.0. Exchange studies at Yale and Hebrew University of Jerusalem.

Skills

Stack: Python (scikit-learn, PydanticAI), Rust, TypeScript/React, Java, SQL, R, C++

Systems and delivery: Temporal, Langfuse, Kubernetes, Jenkins, Grafana/Prometheus