Daniel Fridljand

danielfridljand.de   daniel.fridljand@gmail.com   daniel-fridljand   FridljDa

Summary

AI engineer building production systems in Python and Rust. Ramping into a new technical and stakeholder domain within months. Proudest work to date is an agentic customer-support system I took from discovery to live production as the sole engineer — automating ~87% of B2C traffic since April 2026.

Software Development Experience

Software Consultant - Applied AI — *Munich, Germany*
TNG Technology Consulting — *Dec 2025 – present*

Building and owning production AI systems for enterprise clients, from customer discovery through live operation.

- Architected a hybrid AI-system: deterministic paths for the recurring-intent majority, agentic fallback for the long tail. Shadow-deployed so every response was staff-reviewed before customer contact.
- Built an AI-assisted improvement loop: reviewer feedback became shipped code changes, each gated against a growing Langfuse eval suite.
- Leading development of a privacy-first desktop AI agent (Tauri/React frontend, Rust backend) that lets non-technical staff run agentic workflows over sensitive local files. Hardened for a month until it survived an adversarial 15 person red-team, then rolled out to the internal department.
- Shipped an automated privacy-sensitive review service to Kubernetes. Used by partners to validate customer offers.

Software Consultant - Enterprise Modernization — *Munich, Germany*
TNG Technology Consulting — *Dec 2024 – Dec 2025*

Platform engineer on a multi-team program modernizing a global supply-chain application.

- Deployed a JFrog Artifactory proxy in three days after the work had been repeatedly postponed as too costly, cutting a recurring CI step from eight hours to 30 seconds — saving each developer ~1–2 hours a week.
- Added OWASP dependency scanning to the Jenkins pipeline and built Grafana dashboards over Prometheus metrics, giving stakeholders visibility of CVE reduction throughout the migration.

Data Science Experience

Research Data Analyst - Environmental Epidemiology — *Palo Alto, USA*
Stanford University School of Medicine — *July 2023 – Dec 2023*

Statistical analysis and data tooling for a national study of air pollution and mortality inequality.

- Led the statistical analysis as co-first author of a *Nature Medicine* paper (2024), which found that over half the Black–White age-adjusted mortality gap in the US is attributable to fine-particulate air pollution.
- Built a reproducible pipeline integrating 63M+ death records with pollution estimates and demographic data; open-sourced the analysis code and interactive R Shiny explorer.

Research Data Analyst - Statistical Genomics (Master's Thesis) — *Heidelberg, Germany*
European Molecular Biology Laboratory (EMBL) — *Oct 2021 – May 2022*

Research on multiple-testing methods for large-scale genomics.

- Developed IHW-Forest, a random-forest hypothesis-weighting method with its splitting and weighting logic written in C++ via Rcpp — it recovered over 30% more discoveries than standard corrections.

Education

University of Heidelberg — *Heidelberg, Germany*
M.Sc. in Mathematics — *Oct 2020 – May 2023*

- Grade: 1.0. Exchange studies at Yale and Hebrew University of Jerusalem.

Skills

Stack: Python (PydanticAI, Temporal, FastAPI), Rust, Java, TypeScript/React, C++

Systems and delivery: Kubernetes, ArgoCD, Jenkins, Grafana/Prometheus, ELK, Temporal, Langfuse